

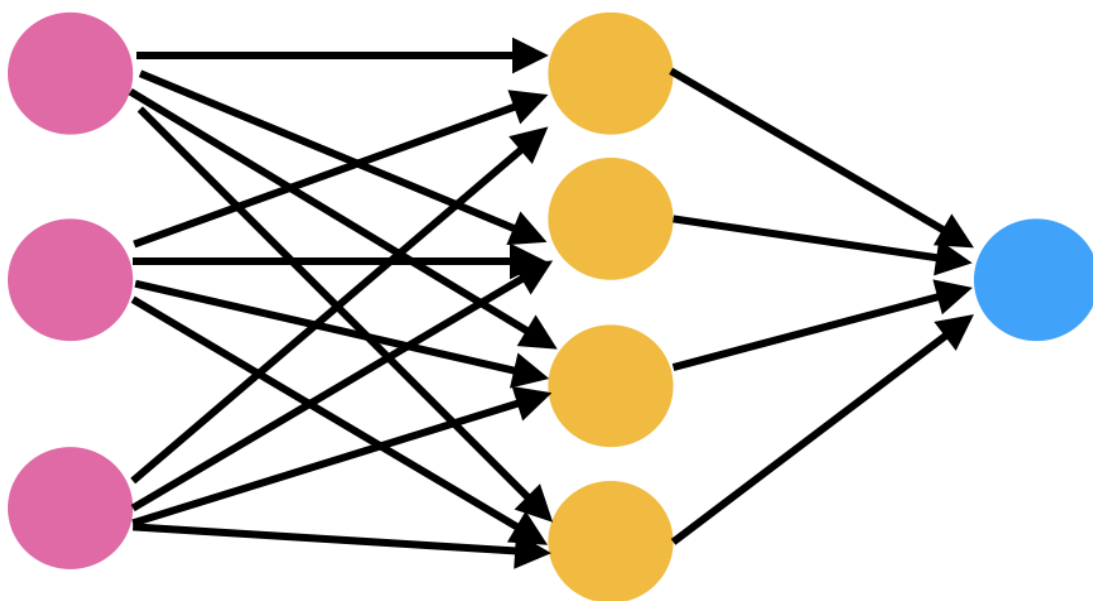
Neural Network Illustration

Neural nets are a means of doing machine learning, in which a computer learns to perform some task by analyzing training examples. Usually, the examples have been hand-labeled in advance. An object recognition system, for instance, might be fed thousands of labeled images of cars, houses, coffee cups, and so on, and it would find visual patterns in the images that consistently correlate with particular labels.

Modeled loosely on the human brain, a neural net consists of thousands or even millions of simple processing nodes that are densely interconnected. Most of today's neural nets are organized into layers of nodes, and they're "feed-forward," meaning that data moves through them in only one direction. An individual node might be connected to several nodes in the layer beneath it, from which it receives data, and several nodes in the layer above it, to which it sends data.

To each of its incoming connections, a node will assign a number known as a "weight." When the network is active, the node receives a different data item — a different number — over each of its connections and multiplies it by the associated weight. It then adds the resulting products together, yielding a single number. If that number is below a threshold value, the node passes no data to the next layer. If the number exceeds the threshold value, the node "fires," which in today's neural nets generally means sending the number — the sum of the weighted inputs — along all its outgoing connections.

When a neural net is being trained, all of its weights and thresholds are initially set to random values. Training data is fed to the bottom layer — the input layer — and it passes through the succeeding layers, getting multiplied and added together in complex ways, until it finally arrives, radically transformed, at the output layer. During training, the weights and thresholds are continually adjusted until training data with the same labels consistently yield similar outputs.



The pink ones are input layer units, for our illustration we have 400 such units and 25 yellow units which are the hidden layer which helps us gain complexity without making the model computationally expensive. For the output the blue layer is the output layer. Since we have 10 numbers to classify from 0 to 9 we have 10 output units for the illustration.

Importing Libraries and Setting options.

```
1 | import numpy as np
2 | import pandas as pd
3 | import matplotlib.pyplot as plt
4 | %matplotlib inline
5 | import seaborn as sns
```

```
1 | pd.set_option('display.max_columns', None)
2 | pd.set_option('display.max_rows', 150)
3 | pd.set_option('display.max_seq_items', None)
4 | sns.set_context('notebook')
5 | sns.set_style('darkgrid')
```

```

1 | from scipy.io import loadmat
2 | traindata = loadmat('data/ex4data1.mat')
3 | print(traindata.keys())
4 | # Let X and y be independent and dependent variable
5 | X = traindata['X']
6 | y = traindata['y']
7 |
8 |
9 | print("Shape of X",X.shape)
10 | print("Shape of y",y.shape)

```

```

1 | dict_keys(['__header__', '__version__', '__globals__', 'X', 'y'])
2 | Shape of X (5000, 400)
3 | Shape of y (5000, 1)

```

The shape of the data tells us that we have 5000 small images of handwritten numbers and their correct answers which shall be used to train our model.

Visualizing the Data

Randomly choosing 9 such numbers and displaying how they look

```

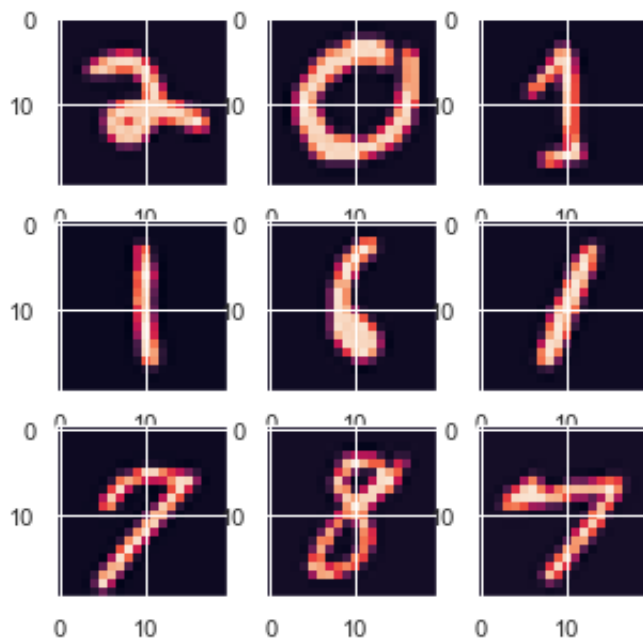
1 | import matplotlib.image as mpimg
2 | def PlotNumbers(X,Nrows=4):
3 |     if Nrows > 100 :
4 |         print("Too Large to Display")
5 |         return
6 |     nrow = int(np.sqrt(Nrows))
7 |     ncol = int(np.sqrt(Nrows))
8 |     m = X.shape[0]
9 |     fig,axis = plt.subplots(nrow,ncol,figsize=(5,5))
10 |    for i in range(ncol):
11 |        for j in range(ncol):
12 |            axis[i,j].imshow(X[np.random.randint(0,m),:].reshape(20

```

```

1 | PlotNumbers(X,9)

```



Model Representation

```

1 | #Introduce Bias
2 | X = np.c_[np.ones(X.shape[0]),X]
3 |
4 | #Loading Weights
5 | datafile = 'data/ex4weights.mat'
6 | mat = loadmat( datafile )
7 | Theta1, Theta2 = mat['Theta1'], mat['Theta2']
8 |
9 | input_layer_size = 400
10 | hidden_layer_size = 25
11 | output_layer_size = 10
12 | n_training_samples = X.shape[0]

```

```

1 | print("Shape of Theta1",Theta1.shape)
2 | print("Shape of Theta2",Theta2.shape)

```

```

1 | Shape of Theta1 (25, 401)
2 | Shape of Theta2 (10, 26)

```

Preparing for Flattening and Reshaping

We will use this many time so lets create a function which will flatten and reshape our parameters (Thetas)

```

1 | # 401x25 + 26x10 = 10285
2 | def FlattenParams(theta1, theta2):
3 |     return np.r_[theta1.ravel(), theta2.ravel()].reshape(((input_layer_size +
4 |         (hidden_layer_size + 1) * output_layer_size), 1))
5 |
6 | # Theta1 25x401 and Theta2 10x26
7 | def ReshapeParams(flattenedTheta):
8 |     Theta1 = flattenedTheta[0:(hidden_layer_size * (input_layer_size + 1))]
9 |     Theta2 = flattenedTheta[(hidden_layer_size * (input_layer_size + 1)):]
10 |    return [Theta1, Theta2]

```

Random Initialization of Parameters

```

1 | def Rand_InitializeTheta(lin, lout):
2 |     epsilon = (6 / (lin + lout)) ** 1/2
3 |     W = np.random.rand(lin, lout + 1) * (2 * epsilon) - epsilon
4 |     return W

```

Cost Function and Forward Propagation

Cost function determines the loss or error that our model is showing with the parameters that we currently have. In order to calculate the loss we have to forward propagate with the available parameters to calculate the hypothesis value. (prediction value) which will be used again, later in the prediction function

```

1  def sigmoid(z):
2      return 1/(1 + np.exp(-1*z))
3
4  def NNCostfunction(thetas,X,y,lamda=0):
5
6      Theta1,Theta2 = ReshapeParams(thetas)
7
8      m = X.shape[0]
9      J = 0
10     y_matrix = pd.get_dummies(y.ravel())
11
12     # Forward Propogation OR Feed Forward
13     a1 = X # 5000x401
14
15     z2 = sigmoid(np.dot(X,Theta1.T)) # 5000x401 * 401x25 = 5000x25
16     a2 = np.c_[np.ones(X.shape[0]),z2] # 5000x26
17
18     z3 = np.dot(a2,Theta2.T) # 5000x10
19     a3 = sigmoid(z3) # 5000x10
20
21     epsilon = 1e-5
22     # If we use vectorized we do not need to sum explicitly
23     J = (-1/m) * np.sum(np.sum(y_matrix * np.log(a3+epsilon) + (1 -
24     #J = (-1/m) * np.sum(np.sum(np.log(a3)*y_matrix + np.log(1-a3)*
25     #J = (-1/m) * np.sum( np.dot(np.log(a3).T,y_matrix) + np.dot(n
26     J = J + (lamda/(2*m)) * (np.sum(np.square(Theta1[:,1:])) + np.s
27
28     return J

```

BackPropagation

Once we have calculated the hypothesis we need to calculate the error or the gap that we have with the predictions and then adjust the parameters going backwards. Starting from the output layer we move backwards to the input layer and find the gradients (partial derivatives of the cost function which is multidimensional space to minimise). We can also calculate the gradients but we will calculate manually and check using the Gradient Check function whether we are close to the numerically calculated derivatives.

```

1  def SigmoidGradient(z):
2      return sigmoid(z)*(1-sigmoid(z))
3
4  def Gradient(thetas,X,y,lamda=0):
5      Theta1,Theta2 = ReshapeParams(thetas)
6
7      m = X.shape[0]
8      J = 0
9      y_matrix = pd.get_dummies(y.ravel())
10
11     # Forward Propogation  OR Feed Forward
12     a1 = X # 5000x401
13
14     z2 = np.dot(X,Theta1.T) # 5000x401 * 401x25 = 5000x25
15     a2 = np.c_[np.ones(X.shape[0]),sigmoid(z2)] # 5000x26
16
17     z3 = np.dot(a2,Theta2.T) # 5000x10
18     a3 = sigmoid(z3) # 5000x10
19
20     ## Backpropogation
21     delta3 = a3 - np.array(y_matrix) #5000x10
22     delta2 = delta3.dot(Theta2[:,1:]) * SigmoidGradient(z2) # 5000x
23     delta1 = None # Since there can be no error with input layer
24
25     Delta2 = delta3.T.dot(a2) # 10x5000* 5000x26 = 10x26
26     Delta1 = delta2.T.dot(a1) # 25x5000 * 5000x401 = 25x401
27
28     # For Both Theta1 and Theta2 do not update the Theta0 i.e Theta
29     # Since REGularization does not affect Theta0.
30
31     Theta1_grad = (1 / m) * Delta1
32     Theta1_grad[:, 1:] = Theta1_grad[:, 1:] + (lamda / m) * Theta1[
33
34     Theta2_grad = (1 / m) * Delta2
35     Theta2_grad[:, 1:] = Theta2_grad[:, 1:] + (lamda / m) * Theta2[
36
37     return np.r_[Theta1_grad.ravel(),Theta2_grad.ravel()]

```

Gradient Checking

```

1 def flattenX(myX):
2     return np.array(myX.flatten()).reshape((n_training_samples*(input_dim+1),n_training_samples))
3
4 def reshapeX(flattenedX):
5     return np.array(flattenedX).reshape((n_training_samples,input_dim+1))
6
7 def checkGradient(theta1,theta2,D1,D2,myX,myy,mylambda=0):
8     myeps = 0.0001
9     flattened = FlattenParams(theta1,theta2)
10    flattenedDs = FlattenParams(D1,D2)
11    myX_flattened = flattenX(myX)
12    n_elems = len(flattened)
13    #Pick ten random elements, compute numerical gradient, compare with BackProp
14    for i in range(10):
15        x = int(np.random.rand()*n_elems)
16        epsvec = np.zeros((n_elems,1))
17        epsvec[x] = myeps
18        cost_high = NNCostfunction(flattened + epsvec,myX,myy,mylambda)
19        cost_low = NNCostfunction(flattened - epsvec,myX,myy,mylambda)
20        mygrad = (cost_high - cost_low) / float(2*myeps)
21        print ("Element: %d. Numerical Gradient = %f. BackProp Gradient = %f" % (x, mygrad, flattenedDs[x]))

```

```

1 thetas = FlattenParams(Theta1,Theta2)
2 print(NNCostfunction(thetas,X,y))
3 grads = Gradient(thetas,X,y)
4 print("Shape of Gradient",grads.shape)
5 D1,D2 = ReshapeParams(grads)

```

```

1 0.28751238099868115
2 Shape of Gradient (10285,)

```

```

1 checkGradient(Theta1,Theta2,D1,D2,X,y)

```

```

1 Element: 4884. Numerical Gradient = -0.000041. BackProp Gradient = -0.000041
2 Element: 8085. Numerical Gradient = -0.000026. BackProp Gradient = -0.000026
3 Element: 8045. Numerical Gradient = 0.000001. BackProp Gradient = 0.000001
4 Element: 4159. Numerical Gradient = -0.000356. BackProp Gradient = -0.000356
5 Element: 2070. Numerical Gradient = 0.000005. BackProp Gradient = 0.000005
6 Element: 7413. Numerical Gradient = 0.000250. BackProp Gradient = 0.000250
7 Element: 6817. Numerical Gradient = 0.000112. BackProp Gradient = 0.000112
8 Element: 10235. Numerical Gradient = -0.000561. BackProp Gradient = -0.000561
9 Element: 5240. Numerical Gradient = 0.000003. BackProp Gradient = 0.000003
10 Element: 9929. Numerical Gradient = -0.000171. BackProp Gradient = -0.000171

```


Prediction and Accuracy

Whichever parameters we have whether it is optimal parameters or randomly initialized parameters, we use it to predict the output and measure the accuracy comparing it with the actual outputs supplied to the function.

```
1 def predictions(thetas,X):
2     Theta1,Theta2 = ReshapeParams(thetas)
3     m = X.shape[0]
4     h1 = sigmoid(np.dot(X,Theta1.T))
5     H1 = np.append(np.ones((m,1)),h1,axis=1)
6     h2 = sigmoid(np.dot(H1,Theta2.T))
7     pred = np.argmax(h2,axis=1)+1
8     return pred
9
10 def accuracy(pred,y):
11     pred.shape = (pred.size,1)
12     return np.sum(pred == y)*100/float(y.shape[0])
```

```
1 pred = predictions(thetas,X)
2 print(NNCostfunction(thetas,X,y))
3 print("Accuracy : {} %".format(accuracy(pred,y)))
```

```
1 0.28751238099868115
2 Accuracy : 10.52 %
```

We have a very low accuracy and out of 5000 samples our model could predict 50 odd numbers correctly. Our next objective is to identify parameters which will minimise the loss (cost function) and improve model performance

Learning Parameters Using Gradient Descent / Scipy / FminCG

In order to improve accuracy, let's use fmin conjugate gradient method. To this method we supply our object : i.e. cost function which needs to be minimised. The gradients which will provide the derivatives. (derivatives provide the slope which tells the method to which direction should the method go ; should it increase the value or decrease)

```

1 | from scipy.optimize import fmin_bfgs, fmin_cg
2 | lamda = 0.00
3 | maxIterations = 200
4 | Learnttheta = fmin_cg(NNCostfunction, thetas, fprime = Gradient,
5 |     args=(X,y,lamda), maxiter = maxIterations, disp = 1)
6 | pred = predictions(Learnttheta,X)
7 | print("Accuracy : {} %".format(accuracy(pred,y)))

```

```

1 | Optimization terminated successfully.
2 |     Current function value: 0.011421
3 |     Iterations: 181
4 |     Function evaluations: 652
5 |     Gradient evaluations: 652
6 | Accuracy : 100.0 %

```

```

1 | print(NNCostfunction(Learnttheta,X,y))

```

```

1 | 0.002391784973557313

```

Bam! We are 100% accurate, that probably means our model is too much dependent on training data and when new data comes it might not perform that well. So let's tune the regularization parameter to penalize the parameters. This will reduce the variance and increase the bias in our model.

Final Trial

Now that we have optimized the cost function, let's randomize the parameters and learn the parameters once again using fmin_CG.

```

1 | Theta1 = Rand_InitializeTheta(25,400)
2 | Theta2 = Rand_InitializeTheta(10,25)
3 | Theta1.shape, Theta2.shape

```

```

1 | ((25, 401), (10, 26))

```

```

1 | thetas = FlattenParams(Theta1, Theta2)
2 | print("Loss : ", NNCostfunction(thetas,X,y))
3 | pred = predictions(thetas,X)
4 | print("Accuracy : {} %".format(accuracy(pred,y)))

```

```
1 | Loss : 6.907571098802985
2 | Accuracy : 10.0 %
```

```
1 | lamda = 0.09
2 | maxIterations = 100
3 | Learnttheta = fmin_cg(NNCostfunction, thetas, fprime = Gradient,
4 |     args=(X,y,lamda), maxiter = maxIterations, disp = 1)
5 | print("Loss : ", NNCostfunction(Learnttheta,X,y))
6 | pred = predictions(Learnttheta,X)
7 | print("Accuracy : {} %".format(accuracy(pred,y)))
```

```
1 | Warning: Maximum number of iterations has been exceeded.
2 |     Current function value: 0.140832
3 |     Iterations: 100
4 |     Function evaluations: 237
5 |     Gradient evaluations: 237
6 | Loss : 0.11346050888543582
7 | Accuracy : 99.26 %
```

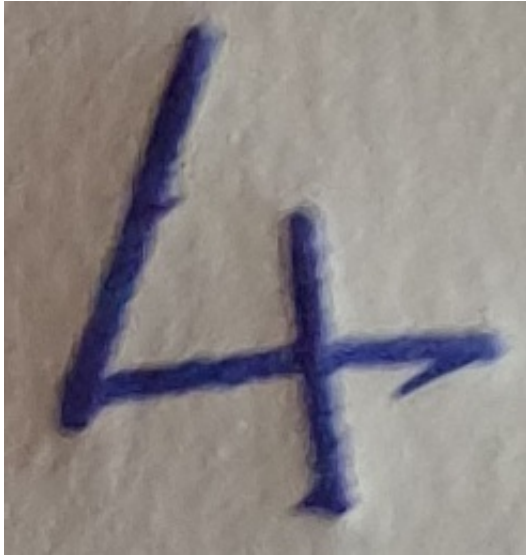
With regularization parameter : lamdas as 0.09 we have achieved 99% accuracy. Time to test the model in real world scenario.

Trying Our Own Image

I wrote number on a paper and imported the image into python, I then converted it into grayscale (because it has two dimensions). An RGB image or colourful image will have 3 dimension for which our model is not currently trained. We need to introduce one more dimension in our theta parameters for our model to be able to read through colourful images. I then scaled the image to a 20x20 : 400 pixel image which suits our model.

```
1 | import cv2
2 | import PIL
3 | from PIL import Image
```

To provide a context, here is the image i clicked



```
1 | basewidth = 20
2 | img = Image.open('number4.jpeg')
3 | wpercent = (basewidth / float(img.size[0]))
4 | #hsize = int((float(img.size[1]) * float(wpercent)))
5 | hsize=basewidth
6 | img = img.resize((basewidth, hsize), PIL.Image.ANTIALIAS)
7 | img.save('number_resized.jpg')
8 |
9 | img = cv2.imread('number_resized.jpg', 0)
10 | img.shape
```

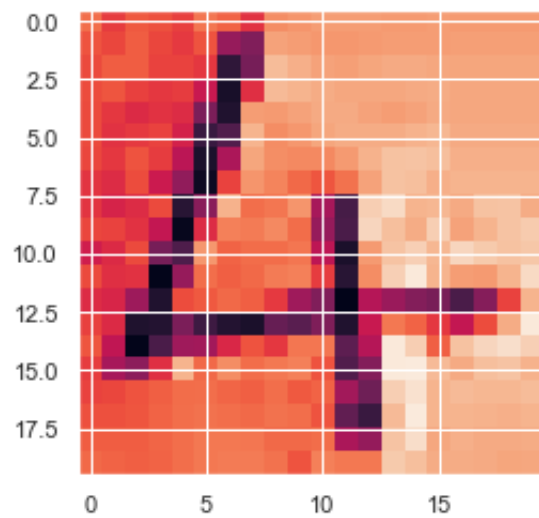
```
1 | (20, 20)
```

```
1 | img = img.ravel()
```

```
1 | img.shape
```

```
1 | (400,)
```

```
1 | plt.imshow(img.reshape(20,20,order="A"))
```



```
1 | img = img.reshape(1,400)
```

```
1 | img.shape
```

```
1 | (1, 400)
```

```
1 | img = np.c_[1,img]
```

```
1 | img.shape
```

```
1 | (1, 401)
```

```
1 | pred = predictions(Learnttheta,img)
2 | print(pred)
```

```
1 | [4]
```

We just supplied one row of data to the model and we use our previously learnt parameters to see if it can correctly predict the image and it did. Our model correctly predicted the number 4.